

Initialization

Run the following code to import the modules you'll need. After you finish the assignment, **remember to run all cells** and save the notebook to your local machine as a PDF for gradescope submission.

```
In [1]: import time
import os
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.patches as patches
```

Download data

In this section we will download the data and setup the paths.

```
In [2]: # Download the data
if not os.path.exists('./content/aerialseq.npy'):
    !wget https://www.cs.cmu.edu/~deva/data/aerialseq.npy -O ./content/aerialseq.npy
if not os.path.exists('./content/antseq.npy'):
    !wget https://www.cs.cmu.edu/~deva/data/antseq.npy -O ./content/antseq.npy
```

Q3: Affine Motion Subtraction

Q3.1: Dominant Motion Estimation (15 points)

```
In [123]: from scipy.interpolate import RectBivariateSpline

def LucasKanadeAffine(It, It1, threshold, num_iters):
    """
    :param It      : (H, W), current image
    :param It1     : (H, W), next image
    :param threshold : (float), if the length of dp < threshold, terminate the
    optimization
    :param num_iters : (int), number of iterations for running the optimization

    :return: M      : (2, 3) The affine transform matrix
    """
    # Initial M
    M = np.array([[1.0, 0.0, 0.0], [0.0, 1.0, 0.0]])
    It_rbs = RectBivariateSpline(np.arange(It.shape[0]), np.arange(It.shape
[1]), It)
    It1_rbs = RectBivariateSpline(np.arange(It1.shape[0]), np.arange(It1.shape
[1]), It1)
    x_grid, y_grid = np.meshgrid(np.arange(It.shape[1]), np.arange(It.shape
[0]))
    old_coords = np.vstack([x_grid.flatten(), y_grid.flatten(), np.ones_like(x
_grid.flatten())])
    # ===== your code here! =====
    for ii in range(num_iters):
        new_coords = M @ old_coords

        mask = ((new_coords[0,:] >= 0) &
                 (new_coords[0,:] < It.shape[1]) &
                 (new_coords[1,:] >= 0) &
                 (new_coords[1,:] < It.shape[0]))

        It1_warp = It1_rbs.ev(new_coords[1], new_coords[0]).reshape(It.shape)

        # dW/dp = [[x, y, 1, 0, 0, 0], [0, 0, 0, x, y, 1]]

        # grad_x = It1_rbs.ev(y_grid, x_grid, dy=1)
        # grad_y = It1_rbs.ev(y_grid, x_grid, dx=1)

        grad_x = np.gradient(It1_warp, axis=1).flatten()
        grad_y = np.gradient(It1_warp, axis=0).flatten()

        # dI/dx' = [dI/dx, dI/dy]
        # dI/dx' * dW/dp = [dI/dx * x, dI/dx * y, dI/dx, dI/dy * x, dI/dy * y,
dI/dy]

        A = np.vstack([grad_x.flatten() * x_grid.flatten(),
                        grad_x.flatten() * y_grid.flatten(),
                        grad_x.flatten(),
                        grad_y.flatten() * x_grid.flatten(),
                        grad_y.flatten() * y_grid.flatten(),
                        grad_y.flatten()]).T
```

```
b = (It - It1_warp).flatten()
dp = np.linalg.lstsq(A[mask], b[mask], rcond=None)[0]

if np.linalg.norm(dp) < threshold:
    break
M += np.reshape(dp, M.shape)

# ===== End of code =====
return M
```

Debug Q3.1

Feel free to use and modify the following snippet to debug your implementation. The snippet simply visualizes the translation resulting from running LK on a single frame. When you warp the source frame using the obtained transformation matrix, it should resemble the target frame.

```

In [124]: import cv2

num_iters = 200
threshold = 0.01
seq = np.load("./content/aerialseq.npy")
It = seq[:, :, 0]
It1 = seq[:, :, 10]


# Source frame
plt.figure()
plt.subplot(1,3,1)
plt.imshow(It, cmap='gray')
plt.title('Source image')

# Target frame
plt.subplot(1,3,2)
plt.imshow(It1, cmap='gray')
plt.title('Target image')

# Warped source frame
M = LucasKanadeAffine(It, It1, threshold, num_iters)
warped_It = cv2.warpAffine(It, M, (It.shape[1], It.shape[0]))
plt.subplot(1,3,3)
plt.imshow(warped_It, cmap='gray')
plt.title('Warped Source image')

```

Out[124]: Text(0.5, 1.0, 'Warped Source image')



Q3.2: Moving Object Detection (10 points)

```

In [156]: import numpy as np
from scipy.ndimage import binary_erosion
from scipy.ndimage import binary_dilation
from scipy.ndimage import affine_transform
import scipy.ndimage
import cv2

def SubtractDominantMotion(It, It1, num_iters, threshold, tolerance):
    """
    :param It      : (H, W), current image
    :param It1     : (H, W), next image
    :param num_iters : (int), number of iterations for running the optimization
    :param threshold : (float), if the length of dp < threshold, terminate the
    optimization
    :param tolerance : (float), binary threshold of intensity difference when
    computing the mask
    :return: mask    : (H, W), the mask of the moved object
    """
    mask = np.ones(It.shape, dtype=bool)

    # ===== your code here! =====
    M = LucasKanadeAffine(It, It1, threshold, num_iters)
    warped_It = cv2.warpAffine(It, -M, It.shape)
    # ===== End of code =====
    mask = np.abs(It1 - warped_It.T) > tolerance
    mask = binary_erosion(mask)
    mask = ~binary_dilation(mask)
    mask[-1,:] = ~mask[-1,:]
    return mask

```

Q3.3: Tracking with affine motion (10 points)

```
In [147]: from tqdm import tqdm

def TrackSequenceAffineMotion(seq, num_iters, threshold, tolerance):
    """
    :param seq      : (H, W, T), sequence of frames
    :param num_iters : int, number of iterations for running the optimization
    :param threshold : float, if the length of dp < threshold, terminate the o
    ptimization
    :param tolerance : (float), binary threshold of intensity difference when
    computing the mask
    :return: masks   : (T, 4) moved objects for each frame
    """
    H, W, N = seq.shape

    rects = []
    It = seq[:, :, 0]
    masks = []

    # ===== your code here! =====
    for i in tqdm(range(1, seq.shape[2])):
        masks.append(SubtractDominantMotion(seq[:, :, i-1], seq[:, :, i], num_iter
s, threshold, tolerance))

    # ===== End of code =====
    masks = np.stack(masks, axis=2)
    return masks
```

Q3.3 (a) - Track Ant Sequence

```
In [148]: seq = np.load("./content/antseq.npy")

# NOTE: feel free to play with these parameters
num_iters = 1000
threshold = 0.01
tolerance = 0.2

tic = time.time()
masks = TrackSequenceAffineMotion(seq, num_iters, threshold, tolerance)
toc = time.time()
print('\nAnt Sequence takes %f seconds' % (toc - tic))



0%|          | 0/124 [00:00<?, ?it/s]100%|██████████| 124/124 [00:16<00:00,
7.38it/s]



Ant Sequence takes 16.857641 seconds
```

```
In [149]: frames_to_save = [29, 59, 89, 119]

# TODO: visualize
for idx in frames_to_save:
    frame = seq[:, :, idx]
    mask = masks[:, :, idx]

    plt.figure()
    plt.imshow(frame, cmap="gray", alpha=0.5)
    plt.imshow(np.ma.masked_where(np.invert(mask), mask), cmap='winter', alpha
=0.8)
    plt.axis('off')
```





Q3.3 (b) - Track Aerial Sequence

```
In [157]: seq = np.load("./content/aerialseq.npy")

# NOTE: feel free to play with these parameters
num_iters = 1000
threshold = 0.01
tolerance = 0.3

tic = time.time()
masks = TrackSequenceAffineMotion(seq, num_iters, threshold, tolerance)
toc = time.time()
print('\nAnt Sequence takes %f seconds' % (toc - tic))
```

```
100%|██████████| 149/149 [00:42<00:00, 3.53it/s]
```

```
Ant Sequence takes 42.298250 seconds
```

```
In [153]: frames_to_save = [29, 59, 89, 119]

# TODO: visualize
for idx in frames_to_save:
    frame = seq[:, :, idx]
    mask = masks[:, :, idx]

    plt.figure()
    plt.imshow(frame, cmap="gray", alpha=0.5)
    plt.imshow(np.ma.masked_where(np.invert(mask), mask), cmap='winter', alpha
=0.8)
    plt.axis('off')
```

